

## Contact

Name: Anya Rajan

Email: [arajan2@ucsc.edu](mailto:arajan2@ucsc.edu)

## Instructions to Run

1. Open terminal
2. cd into the directory that holds the file naive\_bayes.py
3. Once in the correct directory, run the following command:
  - python3 naive\_bayes.py

## Expected Output

The expected output was between 90-95% accuracy.

## My Output

```
[(base) rachelcarlson-100-64-73-162:project2 anyarajan$ python3 naive_bayes.py  
Train Accuracy: 0.9583422804380503  
Test Accuracy: 0.946
```

## Understanding of Why the Program Works

My task was to fill in the code for the “fit” and “predict” functions. After the “fit” function runs, the following had to be true:

1. self.num\_train\_hams is set to the number of ham emails given.
2. self.num\_train\_spams is set to the number of spam emails given.
3. self.word\_counts\_spam is a dictionary where word\_counts\_spam[word] is the number of spam emails which contain this word.

4. `self.word_counts_ham` is a dictionary where `word_counts_ham[word]` is the number of ham emails which contain this word.

To set `num_train_hams` to the number of ham emails given, I found the length of the `train_hams` list passed into the function with `len()`. I did the same, respectively, for `num_train_spams`. To populate the `self.word_counts_ham` dictionary, I used an outer for loop to loop over the ham emails given. Inside this loop, I made a set that holds all the words in the email with `word_set(train_hams[i])`, the counter of the loop being “i”. Then, I began an inner loop that loops over the words in this set. Inside this for loop, I set the element `x` (counter of the loop) of `word_counts_ham` equal to `self.word_counts_ham.get(x, 0) + 1` to keep track of the word and how many times it has been seen in the email. I did the same, respectively, for spam.

My next task was to create the “predict” function. In this, I had to predict whether or not the email passed to the function is spam or ham. I first found the probability of spam emails in the total set by dividing the total number of spam by the sum of total spam and total ham. I also found the probability of ham emails this way. After this, I created the set of words that are in the email passed in with `self.word_set(filename)`. I then looped over the set in a for loop. Inside this for loop, I added to the variable `word_probs_spam`, which I initialized to zero at the beginning of this function. I added the log of the *number of spam emails with the word* (plus one to prevent zero probabilities) divided by the *number of total spam emails* (plus two to prevent zero probabilities) to the `word_probs_spam` variable.

To find the final probability of the email being spam, I added `word_probs_spam` with the log of the total probability of getting a spam email (calculated at the beginning of this function). Normally, you would multiply all of these quantities together, but since we are using the log function, we have to add them. I did the same process for ham emails. After calculating both

final probabilities, I checked if the probability of spam was higher than ham. If it was, the program returns the spam label. Otherwise, it returns the ham label.

\* Using log accounts for underflow problems since we are multiplying many small numbers.

NOTE: For calculating the final probabilities, we are essentially using the LTP, but we do not need to worry about the denominators for both spam and ham because they are the same.

Therefore, we are just calculating and comparing the numerators.